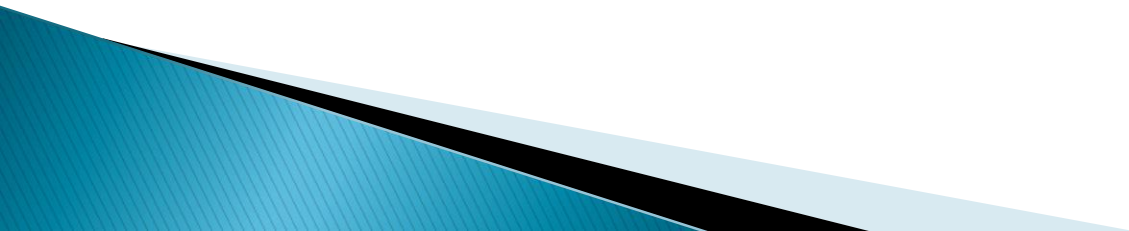



Motion

Optical Flow

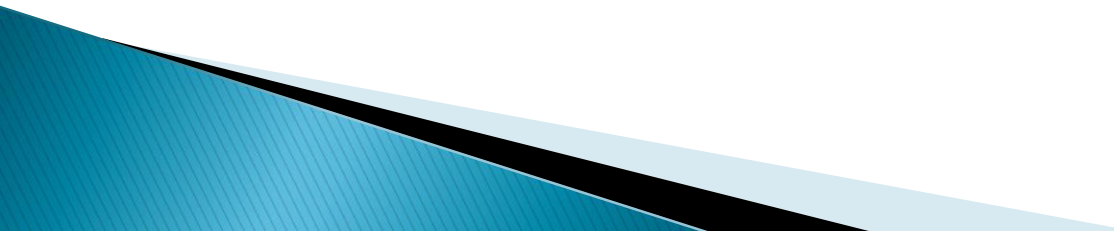


Video

- ▶ A video is a sequence of frames captured over time
- ▶ Now our image data is a function of space (x, y) and time (t)





Uses of motion in computer vision

- ▶ 3D structure reconstruction
 - ▶ Object segmentation from motion
 - ▶ Learning and tracking of dynamical models
 - ▶ Event and activity recognition
 - ▶ Improving video quality
- 

Motion Field


- ▶ The motion field is the projection of the 3D scene motion into the image.



Optical Flow

- ▶ **Definition:** optical flow is the *apparent* motion of brightness patterns in the image.
- ▶ **Note:** apparent motion can be caused by lighting changes without any actual motion
 - Think of a uniform rotating sphere under fixed lighting vs. a stationary sphere under moving illumination


Optical Flow




Picture courtesy of Selim Temizer - Learning and Intelligent Systems (LIS) Group, MIT

Estimating optical flow

- ▶ Given two subsequent frames, estimate the apparent motion field $u(x,y)$ and $v(x,y)$ between them.
- ▶ Key assumptions
 - Brightness constancy: projection of the same point looks the same in every frame
 - Small motion: points do not move very far
 - Spatial coherence: points move like their neighbors



$I(x,y,t-1)$



$I(x,y,t)$

Color/Brightness Constancy





Figure 1.5: Data conservation assumption. The highlighted region in the right image looks roughly the same as the region in the left image, despite the fact that it has moved.

$$I(x, y, t) = I(x + u(x, y), y + v(x, y), t + 1)$$

The brightness constancy constraint



Brightness Constancy Equation:

$$I(x, y, t - 1) = I(x + u, y + v, t)$$

Linearizing the right side using Taylor expansion:

$$I(x, y, t - 1) \approx I(x, y, t) + I_x u + I_y v$$

Hence, $I_x u + I_y v + I_t \approx 0$


The brightness constancy constraint

- ▶ How many equations and unknowns per pixel?
 - One equation per pixel
 - u and v are unknown!

$$I_x u + I_y v + I_t = 0$$


$$\nabla I \cdot (u, v) + I_t = 0$$

The aperture problem



Perceived motion


The aperture problem



Actual motion


slide credit: Kristen Grauman

The barber pole illusion



http://en.wikipedia.org/wiki/Barberpole_illusion

The barber pole illusion



http://en.wikipedia.org/wiki/Barberpole_illusion

Solving the aperture problem

- ▶ How to get more equations for a pixel?
- ▶ Spatial coherence constraint: pretend the pixel's neighbors have the same (u,v)
 - E.g., if we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{matrix}$$

Solving the aperture problem

Prob: we have more equations than unknowns

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{matrix} \longrightarrow \text{minimize } \|Ad - b\|^2$$

Solution: solve least squares problem

- minimum least squares solution given by solution (in d) of:

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 & 2 \times 1 \end{matrix}$$

$$\boxed{\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix}} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$ $A^T b$

- The summations are over all pixels in the $K \times K$ window
- This technique was first proposed by Lucas & Kanade (1981)

Conditions for solvability

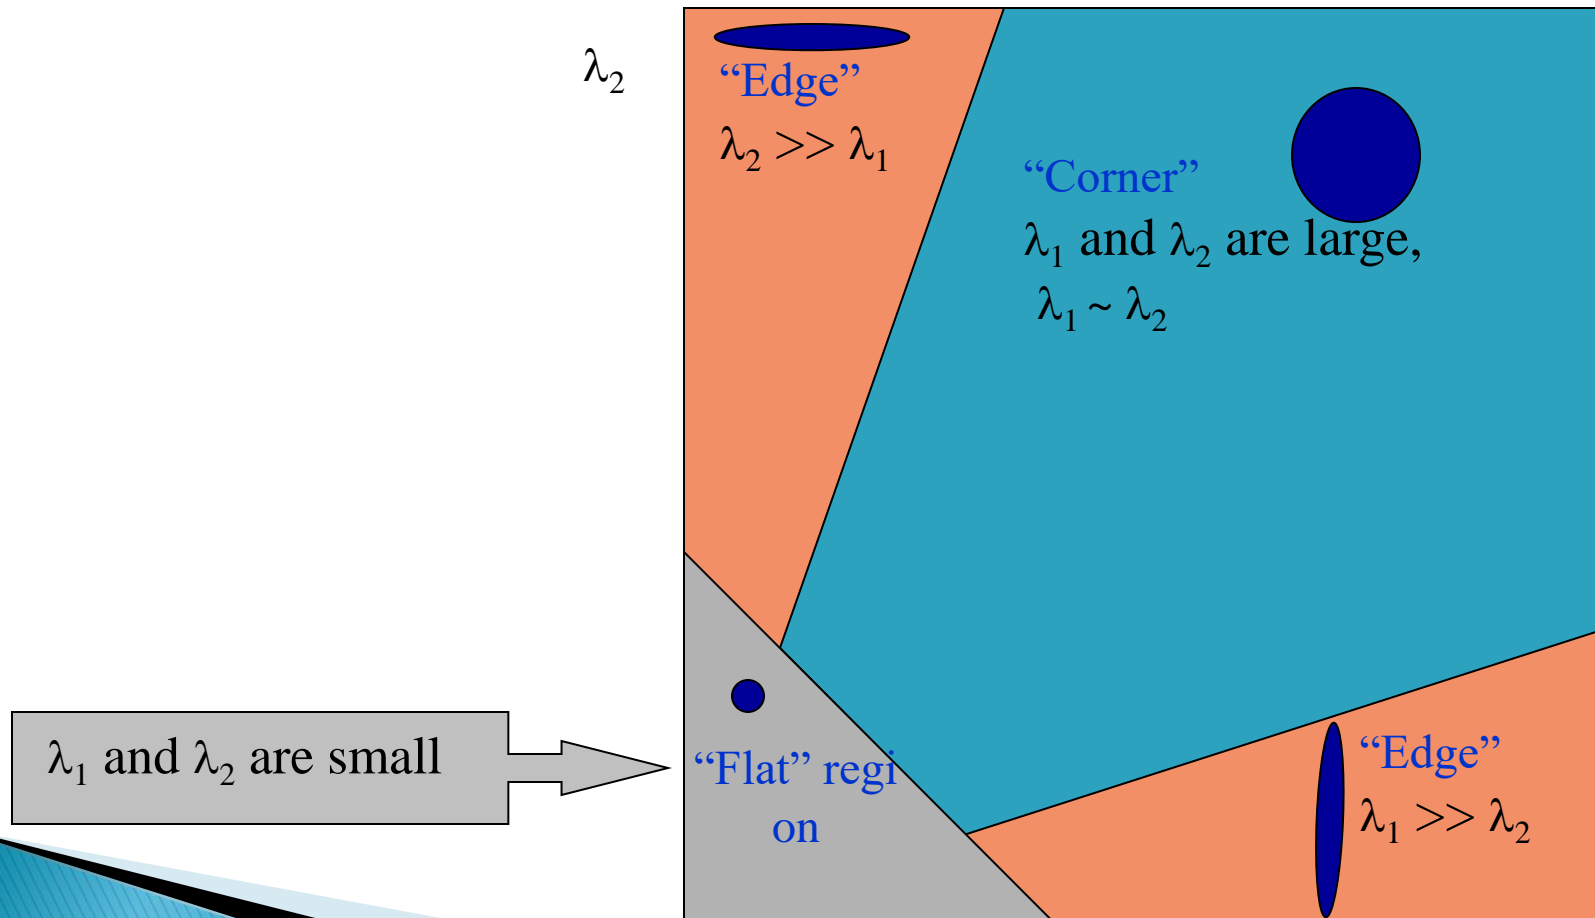
$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A$ $A^T b$

When is this solvable?

- $A^T A$ should be invertible
- $A^T A$ should not be very small
 - eigenvalues λ_1 and λ_2 of $A^T A$ should not be very small
- $A^T A$ should be well-conditioned
 - λ_1 / λ_2 should not be too large (λ_1 = larger eigenvalue)

Recall: Harris Corner Detector...



Edge



- gradients very large or very small
- large λ_1 , small λ_2

Low-texture region




- gradients have small magnitude
- small λ_1 , small λ_2

High-texture region




- gradients are different, large magnitudes
- large λ_1 , large λ_2

Background Subtraction



Video as an “Image Stack”



- ▶ Can look at video data as a spatio-temporal volume
 - If camera is stationary, each line through time corresponds to a single ray in space

Background Subtraction

- ▶ Given an image (mostly likely to be a video frame), we want to identify the **foreground objects** in that image!



Motivation

- ▶ In most cases, objects are of interest, not the scene.
- ▶ Makes our life easier: less processing costs, and less room for error.

Slide credit: Birgi Tamersoy

Background subtraction

- ▶ Simple techniques can do ok with static camera
- ▶ ...But hard to do perfectly
- ▶ Widely used:
 - Traffic monitoring (counting vehicles, detecting & tracking vehicles, pedestrians),
 - Human action recognition (run, walk, jump, squat),
 - Human-computer interaction
 - Object tracking

Simple Approach

Image at time t :

$$I(x, y, t)$$



Background at time t :

$$B(x, y, t)$$



$| > Th$

1. Estimate the background for time t .
2. Subtract the estimated background from the input frame.
3. Apply a threshold, Th , to the absolute difference to get the **foreground mask**.

But, how can we estimate the background?

Frame Differencing

- ▶ Background is estimated to be the previous frame. Background subtraction equation then becomes:

$$B(x, y, t) = I(x, y, t - 1)$$



$$|I(x, y, t) - I(x, y, t - 1)| > Th$$

- ▶ Depending on the object structure, speed, frame rate and global threshold, this approach may or may **not** be useful (usually **not**).



—



| > Th

Frame Differencing

$Th = 25$




$Th = 50$



$Th = 100$



$Th = 200$



Mean Filter

- In this case the background is the mean of the previous n frames:

$$B(x, y, t) = \frac{1}{n} \sum_{i=0}^{n-1} I(x, y, t - i)$$

$$|I(x, y, t) - \frac{1}{n} \sum_{i=0}^{n-1} I(x, y, t - i)| > Th$$

- For $n = 10$:

Estimated Background



Foreground Mask



Median Filter

- Assuming that the background is more likely to appear in a scene, we can use the median of the previous n frames as the background model:

$$B(x, y, t) = \text{median}\{I(x, y, t - i)\}$$

\Downarrow

$$|I(x, y, t) - \text{median}\{I(x, y, t - i)\}| > Th \text{ where } i \in \{0, \dots, n - 1\}.$$

- For $n = 10$:

Estimated Background



Foreground Mask



Slide credit: Birgi Tamersoy

Average/Median Image

